

TCES 420 - Week 8

Jake McKenzie

November 18, 2018

Homework 6

Section 1 8.9 Explain the difference between internal and external fragmentation.

Answer: Internal fragmentation occurs when memory is divided into fixed sized partitions, by which a process has allocated more memory than is required, making less space available overall. This can be fixed by allowing variable sized partitions or by using dynamic memory.

External fragmentation occurs when memory is divided into variable sized partitions, by which a process has allocated more memory than is required, leading to fewer and fewer contiguous blocks of memory. This can be fixed by allowing for compaction, paging and segmentation.

8.11 Given six memory partitions of $300KB$, $600KB$, $350KB$, $200KB$, $750KB$, and $125KB$ (in order), how would the first-fit, best-fit, and worst-fit algorithms place processes of size $115KB$, $500KB$, $358KB$, $200KB$, and $375KB$ (in order)? Rank the algorithms in terms of how efficiently they use memory.

Answer: Let $P_0 = 115$ KB, $P_1 = 500$ KB, $P_2 = 358$ KB, $P_3 = 200$ KB, and $P_4 = 375$ KB.

Worst Fit < First Fit < Best Fit

Table 1: First Fit

Memory	
P_0	300 KB
P_1	600 KB
P_3	350 KB
	200 KB
P_2, P_4	750 KB
	150 KB

Table 2: Best Fit

Memory	
	300 KB
P_1	600 KB
	350 KB
P_3	200 KB
P_2, P_4	750 KB
P_0	150 KB

Table 3: Worst Fit

Memory	
	300 KB
P_2	600 KB
P_3	350 KB
	200 KB
P_0, P_1	750 KB
	150 KB

8.17 Compare paging with segmentation with respect to how much memory the address translation structures require to convert virtual addresses to physical addresses.

Answer: Segmentation will have external fragmentation because the segments can be arbitrary size and may not be able to fit in empty regions of memory with no internal fragmentation because the size of the segment matches the size of the request. The segment table maintains the base address the length of the segment, and segment number (index). The logical address a program uses is some combination of a segment number and some offset. This logical address is added to the base address is what a program uses to navigate memory.

In paging the logical address space can exceed currently available memory and is accomplished by storing portions of the unused program on disk, instead of memory. At first portions of a program are split up on disk, after which these portions are given specific page addresses, which needn't be contiguous for instance for a program τ and a ψ split in portions $\{\tau_0, \tau_1, \tau_2\}$ and $\{\psi_0, \psi_1\}$ on disk the following ordering is valid:

Table 4: Paging

Page Address	Memory
\vdots	\vdots
$0x8$	τ_0
$0x9$	τ_1
$0xA$	ψ_0
$0xB$	τ_2
$0xC$	ψ_1
\vdots	\vdots

In pages, each process in memory has its own page table. Each frame memory holds exactly a flag for whether it is in memory or on disk in the process's table.

Paging allows you to have larger address space without more physical memory, while segmentation does not. The programmer is not aware of pages,

as it is an operation taken care of by the OS, while segmentation is visible and used by the programmer. Pages do not allow for extra security features while segmentation allows you to. Changing data in segmentation allows for individual compilation of that changed segment while paging requires a complete compile.

8.20 Assuming a 1-KB page size, what are the page numbers and offsets for the following address references (provided as decimal numbers):

Answer: The logical address length is $\log_2 1024 = 10$ and the page number is the rest of the bits. I like working in Hex, denoted by: page number.offset

- a) $3085_{10} = 0x3.0xD = 3_{10}.13_{10}$
- b) $42095_{10} = 0x29.0x6F = 41_{10}.111_{10}$
- c) $215201_{10} = 0xD2.0xA1 = 210_{10}.161_{10}$
- d) $650000_{10} = 0x27A.0x310 = 634_{10}.784_{10}$
- e) $2000001_{10} = 0x200.0x1 = 512_{10}.1_{10}$

8.21 a) The BTV operating system has a 21-bit virtual address, yet on certain embedded devices, it has only a 16-bit physical address. It also has a 2-KB page size. How many entries are there in a conventional single-level page table?

Answer: $2^{10} = 1024$

8.22 What is the maximum amount of physical memory?

Answer: The maximum amount of memory in a BTV operating system is 64 KB.

8.23 Consider a logical address space of 256 pages with a 4-KB page size,

mapped onto a physical memory of 64 frames. How many bits are required in the logical address? How many bits are required in the physical address?

Answer: $4 \text{ KB} = 4 * 2^{10} = 2^{12}$

The number of bits required for logical address is $\log_2 256 + \log_2 2^{12} = \log_2 2^8 + \log_2 2^{12} = 8 + 12 = 20$

The number of bits required for logical address is $\log_2 64 + \log_2 2^{12} = \log_2 2^6 + \log_2 2^{12} = 6 + 12 = 18$

8.24 Consider a computer system with a 32-bit logical address and 4-KB page size. The system supports up to 512 MB of physical memory. How many entries are there in each of the following?

a) a conventional single-level page table:

Answer: Entries = $\frac{\text{logical memory space}}{\text{physical memory}} = \frac{2^{32}}{2^{12}} = 2^{20} = 1\text{MB}$

b) An inverted page table

Answer: Entries = $\frac{\text{physical memory size}}{\text{page size}} = \frac{2^9 * 2^{20}}{2^2 * 2^{10}} = 2^{17} = 128\text{KB}$

8.25 with TLB hitrate = 95%...Consider a paging system with the page table stored in memory:

a) If a memory reference takes 50 nanoseconds, how long does a paged memory reference take?

Answer: $4 * (\text{Time to access page table} \times \text{Time to access word in memory}) = 4 * (50\text{ns} + 50\text{ns}) = 4 * 10\text{ns} = 40\text{ns}$

b) If we add TLBs, and 75 percent of all page-table references are found in the TLBs, what is the effective memory reference time? (Assume that finding a page-table entry in the TLBs takes 2 nanoseconds, if the entry is

present.)

Answer: $0.95 * 200 + 0.05 * 400 = 210\text{ns}$

8.28 Consider the following segment table:

<u>Segment</u>	<u>Base</u>	<u>Length</u>
0	219	600
1	2300	14
2	90	100
3	1327	580
4	1952	96

What are the physical addresses for the following logical addresses?

Answer: a) $219 + 430 = 649$

b) $2300 + 10 = 2310$

c) Trap thrown by OS

d) $1327 + 400 = 1727$

e) Trap thrown by OS

8.29 What is the purpose of paging the page tables?

Answer: Firstly it puts the user page tables in a pageable segment of the system's address space. Secondly it allows the system segment containing the user page tables to be paged. Lastly it allows you to reduce the physical memory requirements of page tables by mapping a portion of the address space to that is actually being used.

8.32 Consider the Intel address-translation scheme shown in Figure 8.22.

a) Describe all the steps taken by the Intel Pentium in translating a logi-

cal address into a physical address.

Answer: The selector is an index into the segment descriptor table. The segment descriptor result plus the original offset is used to produce a linear address with $\text{dir} \rightarrow \text{page} \rightarrow \text{offset}$. The dir is an index into a page directory; the entry from this directory selects the page table, and the page field is an index into that page table. The entry from the page table, plus the offset, is the physical address.

b) What are the advantages to the operating system of hardware that provides such complicated memory translation?

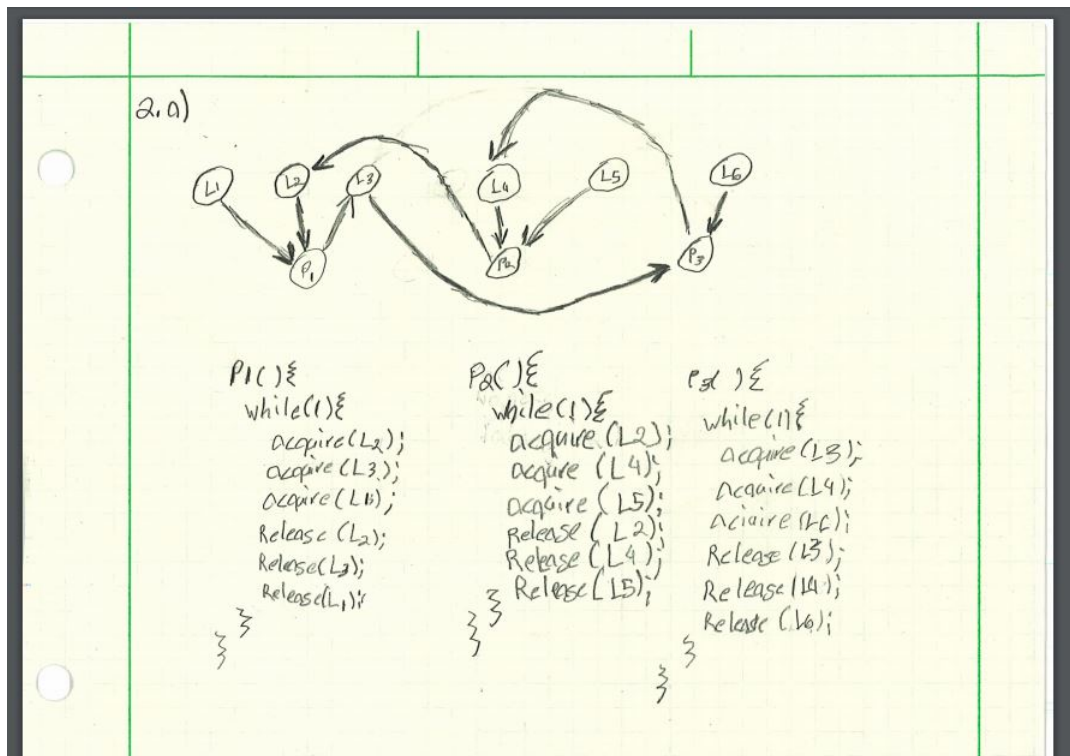
Answer: This allows for a level of abstraction that allows for various hardware configurations. Because it allows for different hardware configurations it allows for more efficiencies. c) Are there any disadvantages to this address-translation system? If so, what are they? If not, why is this scheme not used by every manufacturer?

The restriction here is on how long it takes for table lookup, which while be $O(1)$ time runtime in reality can add up if not taken into account. Caches can help, but you can still have a cache miss.

Section 2 Three processes need to acquire six locks labeled L_1 to L_6 , as shown below(2 points)

a) Use an allocation graph to show the possibility of a deadlock in this code.

Answer:



b) Change the order of some of the acquires to prevent the possibility of any deadlock. You cannot move acquire requests across processes, only change the order inside each process. Use an allocation graph to justify your answer.

Answer: To prevent deadlock we need to acquire resources either in ascending or descending order to prevent deadlock. NOT all cycles make the RAG in deadlock, but if there are no cycles the processes are definitely NOT in deadlock. If only one instance per resource type, then there is a deadlock. If there are several instances per resource type, there is a high likelihood of a deadlock. The probabilistic view tells us to not worry about this problem because it is so rare. UNIX does not take care of this problem because they hold this view.

Section 3 Answer: First P_3 completes making the new available vector $\langle 2, 1, 0, 0 \rangle + \langle 0, 1, 2, 0 \rangle = \langle 2, 2, 2, 0 \rangle$. After which P_2 completes making the available vector $\langle 2, 2, 2, 0 \rangle + \langle 2, 0, 0, 1 \rangle = \langle 4, 2, 2, 1 \rangle$ and

after which P_1 completes making the final availability vector $\langle 4, 2, 2, 1 \rangle + \langle 0, 0, 1, 0 \rangle = \langle 4, 2, 3, 1 \rangle$. So the sequence is $P_3 \rightarrow P_2 \rightarrow P_1$ and the system is not in a deadlock.

Section 4 Answer: R_{01} arrives at 1 and finishes at 5 while R_{22} arrives at 2 and finishes at 6. R_{34} arrives at 3 and finishes at 7. W_{16} arrives at 4 but cannot start until 7 and finishes at 15. R_{47} arrives at 5 but waits until W_{16} completes. R_{49} arrives at 6 but must wait until W_{16} completes. W_{214} arrives at 7 and waits until W_{16} completes, when it does it starts and completes at 23. After this R_{47} and R_{49} are free to go where they both complete at 27.