TCES 420 - Week 8

Jake McKenzie

November 26, 2018

Homework 7

Section 1 9.18 A certain computer provides its users with a virtual memory space of 2^{32} bytes. The computer has 2^{22} bytes of physical memory. The virtual memory is implemented by paging, and the page size is $4,096$ bytes. A user process generates the virtual address 11123456. Explain how the system establishes the corresponding physical location. Distinguish between software and hardware operations.

Answer: If we convert the virtual address from decimal to hex we obtain $11123456_{10} = $ 0x00A9BB00. We can obtain the page size of $2^{(}12)$ thus the first 12 bits are for displacement within pages and we can deduce that the displacement within the table is 20 bits because the table size is $2^{(}20)$. We obtain a 0x00A9B for the pages table displacement and 0xB00 for pages displacement.

9.19 Assume that we have a demand-paged memory. The page table is held in registers. It takes 8 milliseconds to service a page fault if an empty frame is available or if the replaced page is not modified and 20 milliseconds if the replaced page is modified. Memory-access time is 100 nanoseconds. Assume that the page to be replaced is modified 70 percent of the time. What is the maximum acceptable page-fault rate for an effective access time of no more than 200 nanoseconds?

```
syms p
us = 10^-6

us = 1.0000e-06

ns = 10^-9

ns = 1.0000e-09

ms = 10^-3

ms = 1.0000e-03

eqn = 200.*ns == (1-p)*100.*ns + 0.3*8.*ms*p+0.7*20.*ms*p

eqn =
```

$$\frac{7555786372591433}{37778931862957161709568} = \frac{163999\,p}{10000000} + \frac{1}{10000000}$$

```
vpa(solve(eqn,p),5)
```

$$\text{ans} = 6.0976\ 10^{-6}$$

Answer: $p = 6.0976 * 10^{-6}$

When a page fault occurs, the process requesting the page must block while waiting for the page to be brought from disk into physical memory. Assume that there exists a process with five user-level threads and that the mapping of user threads to kernel threads is one to one. If one user thread incurs a page fault while accessing its stack, would the other user threads belonging to the same process also be affected by the page fault—that is, would they also have to wait for the faulting page to be brought into memory? Explain.

Answer: If you have multiple threads, while one thread is blocking which is the case in this example, the other threads in this process will be uneffected by the blocking. Only on that single thread who is blocking is there an negative effect from the page fault.

9.22 The page table shown in Figure 9.32 is for a system with 16-bit virtual and physical addresses and with 4,096-byte pages. The reference bit is

2

set to 1 when the page has been referenced. Periodically, a thread zeroes out all values of the reference bit. A dash for a page frame indicates the page is not in memory. The page-replacement algorithm is localized LRU, and all numbers are provided in decimal.

a) Convert the following virtual addresses (in hexadecimal) to the equivalent physical addresses. You may provide answers in either hexadecimal or decimal. Also set the reference bit for the appropriate entry in the page table.

Answer:
Virtual Address $\sim$ physical address
0xE12C $\sim$ 0x312C
0x3A9D $\sim$ 0xAA9D
0xA9D9 $\sim$ 0x59D9
0x7001 $\sim$ 0xF001
0xACA1 $\sim$ 0x5CA1

b) Using the above addresses as a guide, provide an example of a logical address (in hexadecimal) that results in a page fault.

Answer: 0x1234, page 1 is occupied so the resulting store will result in a fault.

c) From what set of page frames will the LRU page-replacement algorithm choose in resolving a page fault?

Answer: 0x4, 0x8, 0xC and 0xD are all page frames that will be picked by the page replacement algorithm page faults.

9.23 Assume that you are monitoring the rate at which the pointer in the clock algorithm moves. (The pointer indicates the candidate page for replacement.) What can you say about the system if you notice the following behavior:
a) pointer is moving fast.
Answer: If the pointer is moving quickly that would imply there are very few page faults and that we are accessing a lot of pages.

b) pointer is moving slow.

Answer: If the power is moving slowly that would imply there are very many page faults and that we are not accessing many pages.

9.27 Consider a demand-paging system with the following time-measured utilizations:

CPU utilization	20%
Paging disk	97.7%
Other I/O devices	5%

For each of the following, indicate whether it will (or is likely to) improve CPU utilization. Explain your answers.

a. Install a faster CPU.

Answer: A faster CPU, in theory, should reduce its utilization by accomplishing more tasks more quickly.

b. Install a bigger paging disk.

Answer: Having more room for pages will not do anything to CPU usuage.

c. Increase the degree of multiprogramming.

Answer: This will make the CPU utilization worse because the number of page faults will increase as the degree of multiprogramming increases.

d. Decrease the degree of multiprogramming.

Answer: This will make the CPU utilization better because the number of page faults will decrease as the degree of multiprogramming decreases.

e. Install more main memory.

Answer: More main memory means more room for active pages meaning less page faults, leading to less CPU utilization.

f. Install a faster hard disk or multiple controllers with multiple hard disks.

Answer: This will increase the overall throughput of the system, but

the overall utilization statistics should stay constant.

g. Add prepaging to the page-fetch algorithms.

Answer: This will increase cpu utilization as the cpu has to do more work for each paging request.

h. Increase the page size.

Answer: Increasing the size of each page will decrease the amount of overall pages which means there are less pages to fault, reducing cpu utilization.

9.30 A page-replacement algorithm should minimize the number of page faults. We can achieve this minimization by distributing heavily used pages evenly over all of memory, rather than having them compete for a small number of page frames. We can associate with each page frame a counter of the number of pages associated with that frame. Then, to replace a page, we can search for the page frame with the smallest counter.

a) Define a page-replacement algorithm using this basic idea. Specifically address these problems:
i. What is the initial value of the counters?
Answer: counter = 0
ii. When are counters increased?
Answer: Whenever a new page is associated with a frame in current context the counter is increased.
iii. When are counters decreased?
Answer: Whenever a new page is associated with a frame that is not in current context the counter is increased.
iv. How is the page to be replaced selected?
Answer: Whichever page exists in a frame that has the smaller counter where ties are broken by whoever arrives first (FIFO).
b) How many page faults occur for your algorithm for the followingn reference string with four page frames?
$$1, 2, 3, 4, 5, 3, 4, 1, 6, 7, 8, 7, 8, 9, 7, 8, 9, 5, 4, 5, 4, 2$$
Answer:
$$\text{F,F,F,F,F,F,-,F,F,F,F,-,-,F,-,-,-,F,F,-,-,F}$$

I found there were 14 page faults.

c) What is the minimum number of page faults for an optimal pagereplacement strategy for the reference string in part b with four page frames?

$$1, 2, 3, 4, 5, 3, 4, 1, 6, 7, 8, 7, 8, 9, 7, 8, 9, 5, 4, 5, 4, 2$$

Answer:
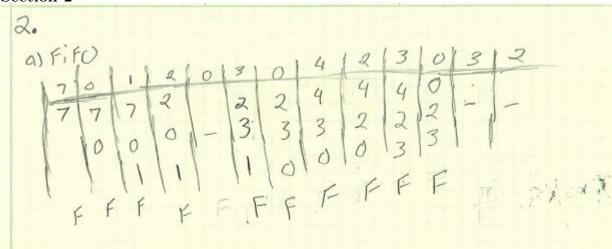
$$F,F,F,F,F,-,-,-,F,F,F,-,-,F,-,-,-,-,F,F$$

If found there were 11 page faults.

9.32 What is the cause of thrashing? How does the system detect thrashing? Once it detects thrashing, what can the system do to eliminate this problem?

Answer:

Thrashing occurs when a computer's virtual memory resources are overused, leading to a constant state of paging and page faults, inhibiting most application-level processing. By tracking working sets, we can detect thrashing by comparing the total working set size to the number of frames of available memory. We can also make use of the size of the working sets when determining which processes to suspend.

Section 2

b) LRU

c) clock

d) optimal

2. e)
For 2. a) I found there were 7 total page faults with miss rate of 70% and
for 2. b) I found there were 6 total page faults with miss rate of 60% and for
2. c) I found there were 5 total page faults with miss rate of 50% and for 2.
d) I found there were 4 total page faults with miss rate of 40%.

Section 3

a) There are 2^5 total pages that take up 2 KBytes each. That gives us logical page number of 5 and an offset of 11 because 2 KBytes $= 2^1 * 2^{10} = 2^{11}$.

b) The width $= \frac{1MB}{2Kb} = \frac{2^{20}}{2^{11}} = 2^9$ so the width is 9 bits. There are 32 pages so there must be a length of 32 to account for each page.

c) Making physical memory space reduce by half means shifting to the left by 1 or to show my work $\frac{2^{19}}{2^{11}} = 2^8$. Now the width is effectively 8 bits.

Section 4

```
1    Hit = 0.98

     Hit = 0.9800

2    MAT = 50*10^-9

     MAT = 5.0000e-08

3    TLB = 2*10^-9

     TLB = 2.0000e-09

4    p = 2*10^-6

     p = 2.0000e-06

5    disk = 5*10^-3

     disk = 0.0050

6    cache = 10*10^-9

     cache = 1.0000e-08

7    HR = 0.9

     HR = 0.9000

8    miss = 200*10^-9

     miss = 2.0000e-07

9    cache_at90 = 0.9*(cache)+0.1*(cache+miss)

     cache_at90 = 3.0000e-08

10   %with cache
11   (1-p)*(Hit*(TLB + cache_at90)+(1-Hit)*(TLB + cache_at90 + cache + cache)) + (p)*(TLB + cache_at90+ disk + cache_at09 + cache_at09)

     ans = 4.2400e-08

12   %without cache
13   (1-p)*(Hit*(TLB+ MAT)+(1-Hit)*(TLB + MAT + MAT)) + (p)*(TLB + MAT + disk)

     ans = 6.3000e-08
```

With a cache I got an Overall EAT $= 42.4$ns and without cache an Overall EAT $= 63$ns.

Section 5

For Semaphore Z there can be at most 10 $(20 - 10 = 10)$ blocked processes because it will decrement the first 10 then can block when $Z = 0$. After which Semaphore Y can there can also be at most 10 $(30 - 20 = 10)$ blocked processes because it will decrement the first 20 then block when $Y = 0$. For

Semaphore X there can be an infinite amount of processes blocked as X can reach zero and stay at zero.